

---

# Cloud Infrastructure

---

# Existing cloud infrastructure

---

- The cloud computing infrastructure at Amazon, Google, and Microsoft (as of mid 2012).
  - Amazon is a pioneer in Infrastructure-as-a-Service (IaaS).
  - Google's efforts are focused on Software-as-a-Service (SaaS) and Platform-as-a-Service (PaaS).
  - Microsoft is involved in PaaS.
- Private clouds are an alternative to public clouds. Open-source cloud computing platforms such as:
  - Eucalyptus,
  - OpenNebula,
  - Nimbus,
  - OpenStack

can be used as a control infrastructure for a private cloud.

# Amazon Web Services (AWS)

---

- AWS → IaaS cloud computing services launched in 2006.
- Businesses in 200 countries used AWS in 2012.
- The infrastructure consists of compute and storage servers interconnected by high-speed networks and supports a set of services.
- An application developer:
  - Installs applications on a platform of his/her choice.
  - Manages resources allocated by Amazon.

# AWS regions and availability zones

- Amazon offers cloud services through a network of data centers on several continents.
- In each *region* there are several availability zones interconnected by high-speed networks.
- An *availability zone* is a data center consisting of a large number of servers.

Region	Location	Availability zones	Cost
US West	Oregon	us-west-2a/2b/2c	Low
US West	North California	us-west-1a/1b/1c	High
US East	North Virginia	us-east-1a/2a/3a/4a	Low
Europe	Ireland	eu-west-1a/1b/1c	Medium
South America	Sao Paulo, Brazil	sa-east-1a/1b	Very high
Asia Pacific	Tokyo, Japan	ap-northeast-1a/1b	High
Asia Pacific	Singapore	ap-southeast-1a/1b	Medium

- Regions do not share resources and communicate through the Internet.



# AWS instances

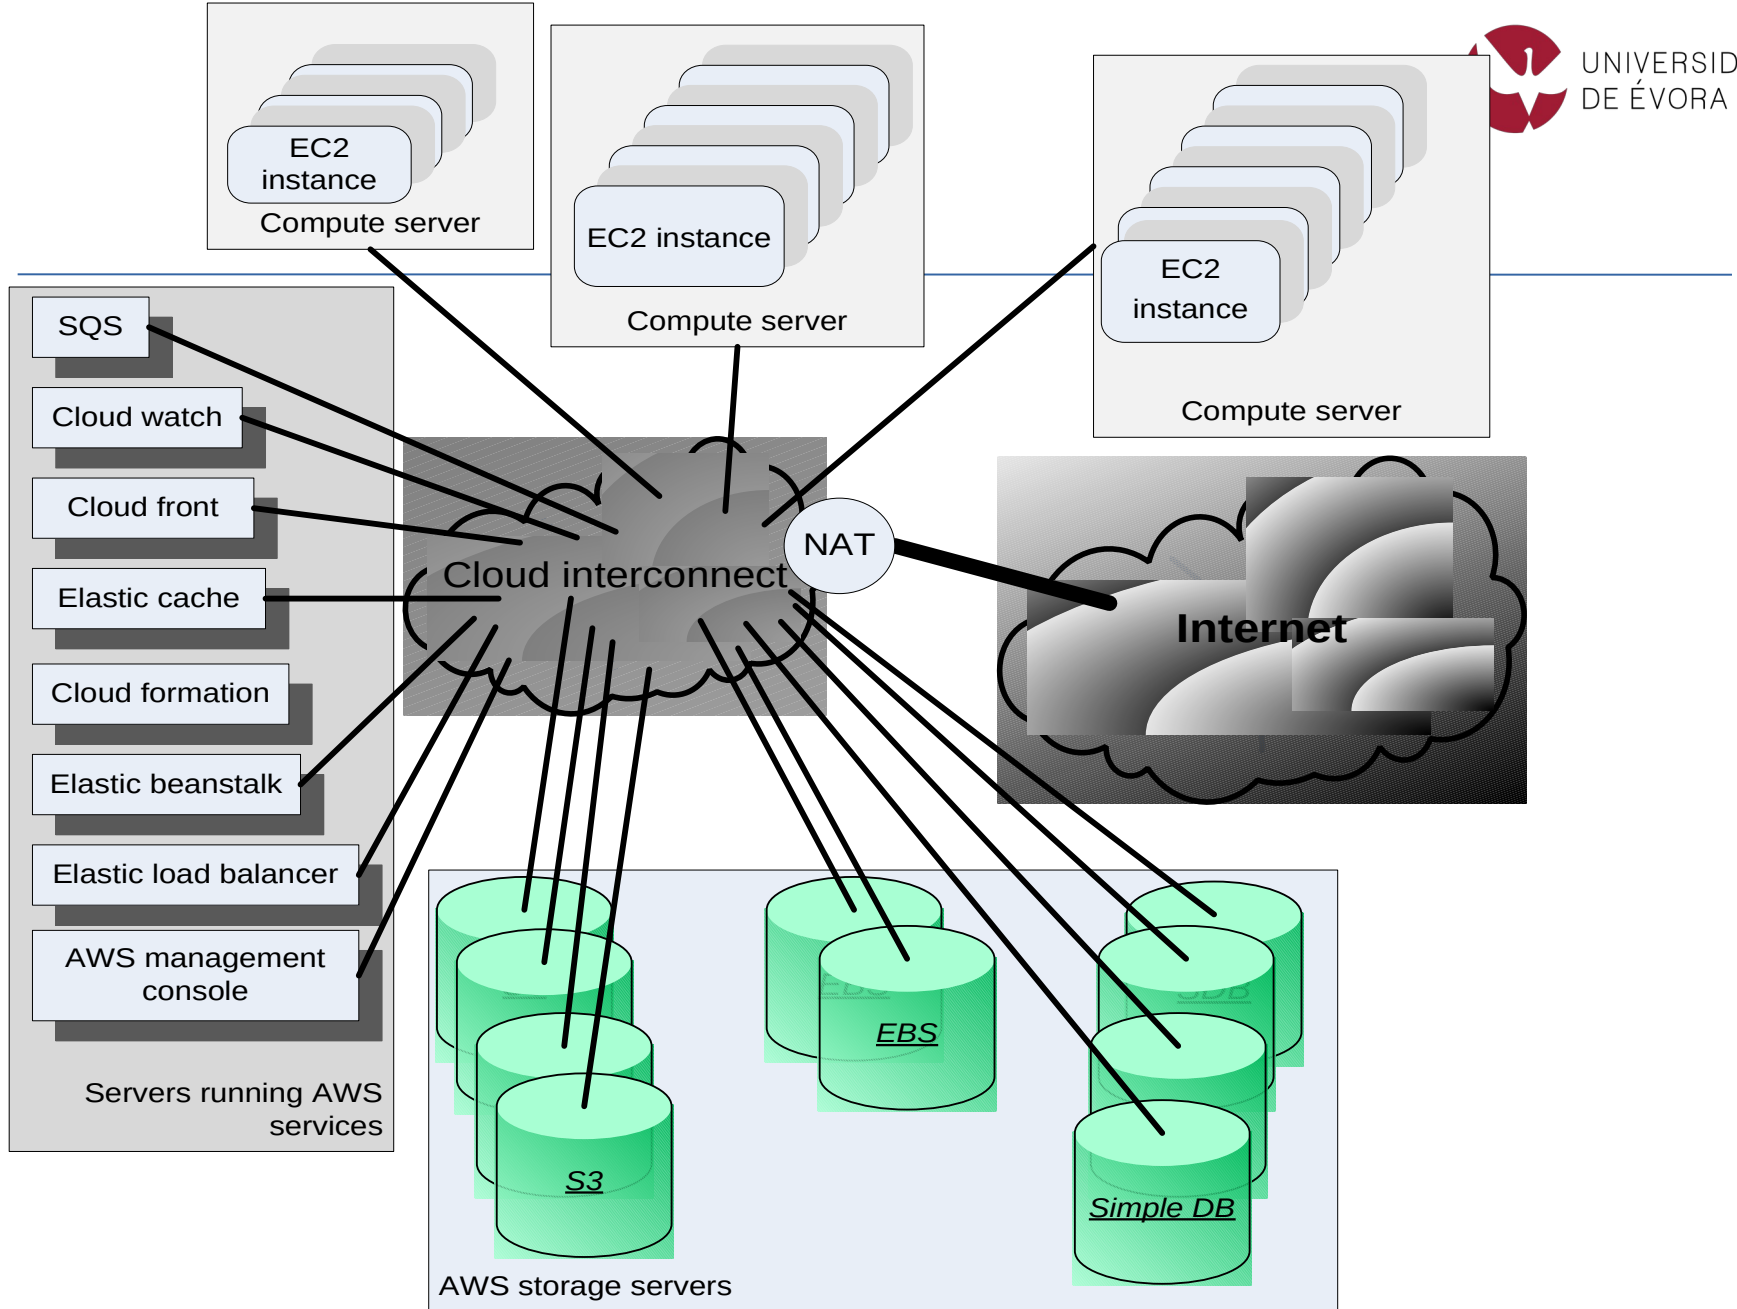
---

- An instance is a virtual server with a well specified set of resources including: CPU cycles, main memory, secondary storage, communication and I/O bandwidth.
- The user chooses:
  - The region and the availability zone where this virtual server should be placed.
  - An instance type from a limited menu of instance types.
- When launched, an instance is provided with a DNS name; this name maps to a
  - *private IP address* → for internal communication within the internal EC2 communication network.
  - *public IP address* → for communication outside the internal Amazon network, e.g., for communication with the user that launched the instance.

# AWS instances (cont'd)

---

- Network Address Translation (NAT) maps external IP addresses to internal ones.
- The public IP address is assigned for the lifetime of an instance.
- An instance can request an *elastic IP address*, rather than a public IP address. The elastic IP address is a static public IP address allocated to an instance from the available pool of the availability zone.
- An elastic IP address is not released when the instance is stopped or terminated and must be released when no longer needed.





# Steps to run an application

---

- Retrieve the user input from the front-end.
- Retrieve the disk image of a VM (Virtual Machine) from a repository.
- Locate a system and requests the VMM (Virtual Machine Monitor) running on that system to setup a VM.
- Invoke the Dynamic Host Configuration Protocol (DHCP) and the IP bridging software to set up MAC and IP addresses for the VM.

# User interactions with AWS

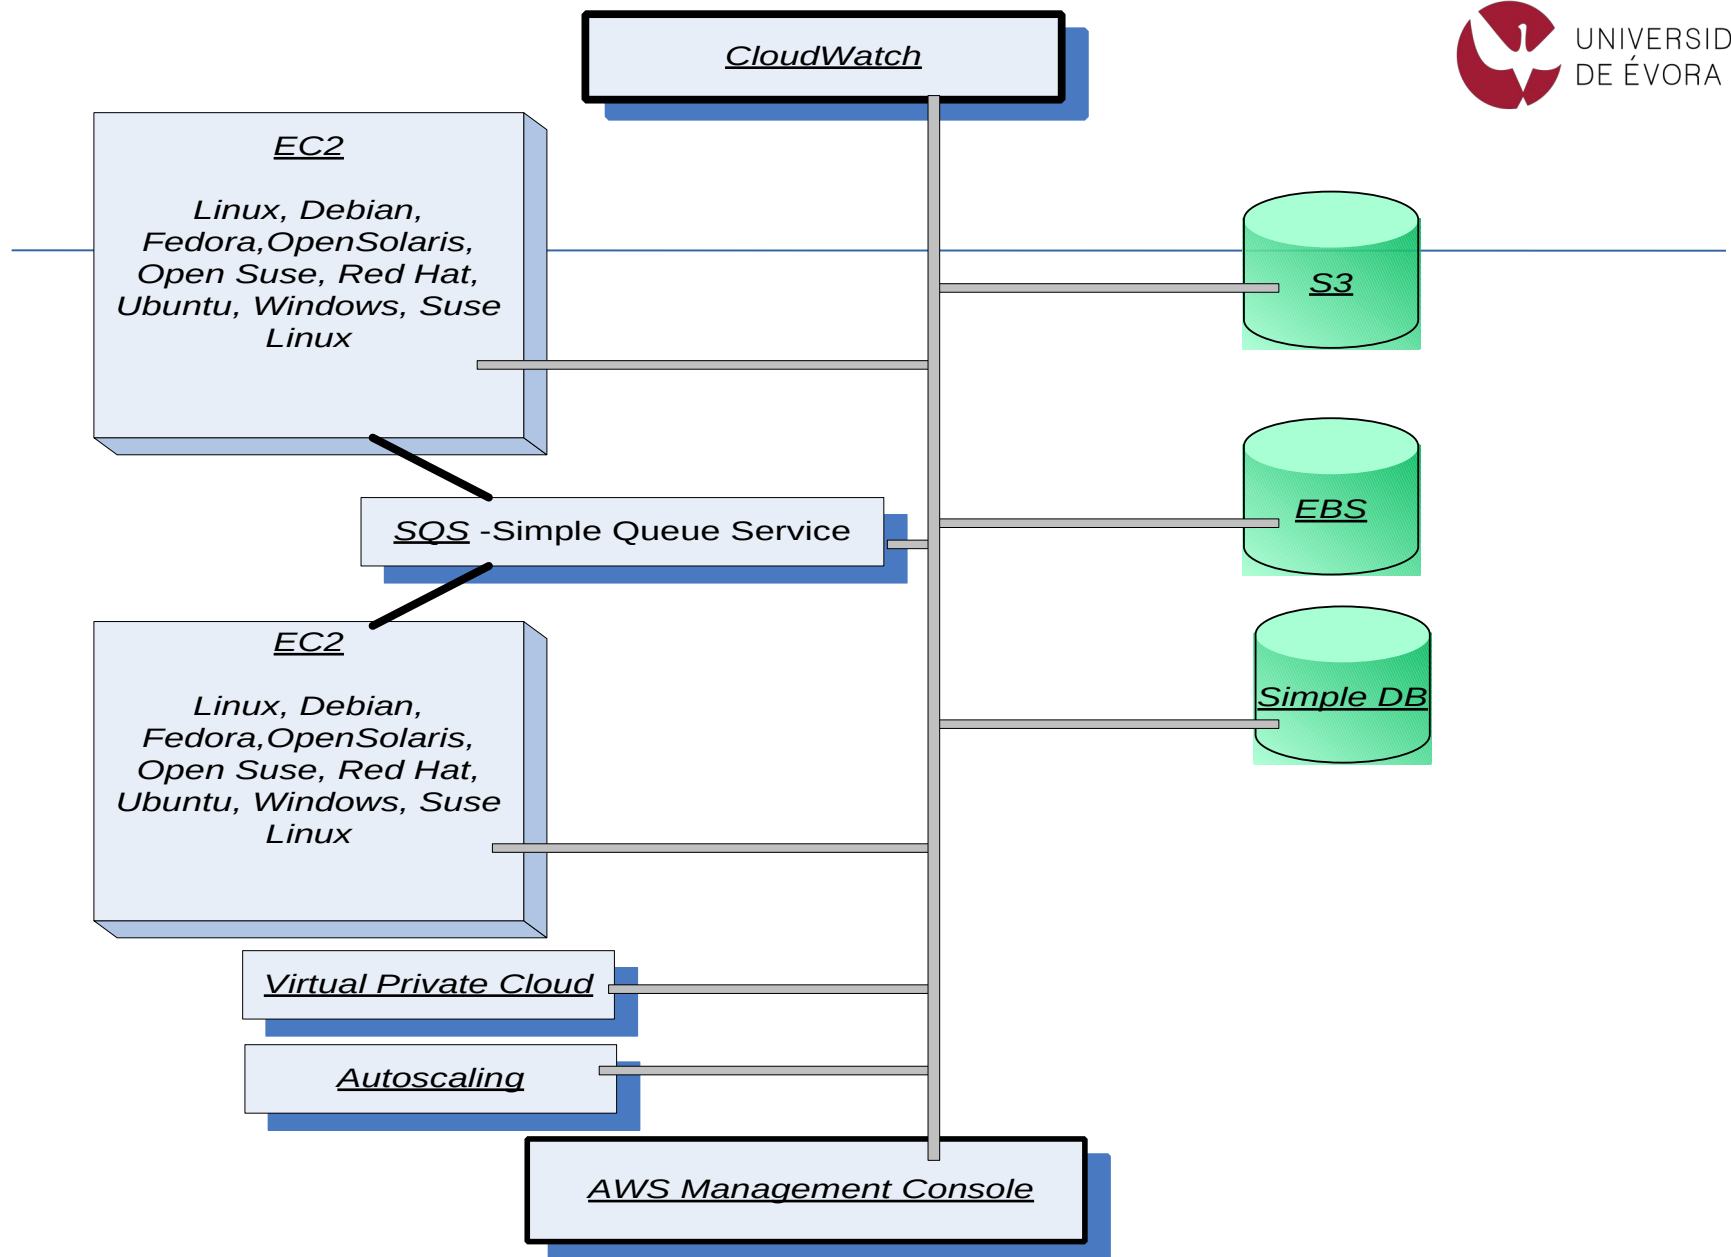
---

- The AWS Management Console. The easiest way to access all services, but not all options may be available.
- AWS SDK libraries and toolkits are provided for several programming languages including Java, PHP, C#, and Objective-C.
- Raw REST requests.

# Examples of Amazon Web Services

---

- *AWS Management Console* - allows users to access the services offered by AWS .
- *Elastic Cloud Computing (EC2)* - allows a user to launch a variety of operating systems.
- *Simple Queuing Service (SQS)* - allows multiple *EC2* instances to communicate with one another.
- *Simple Storage Service (S3), Simple DB, and Elastic Block Storage (EBS)* - storage services.
- *Cloud Watch* - supports performance monitoring.
- *Auto Scaling* - supports elastic resource management.
- *Virtual Private Cloud* - allows direct migration of parallel applications.



# EC2 – Elastic Cloud Computing

---

- *EC2* - web service for launching instances of an application under several operating systems, such as:
  - Several Linux distributions.
  - Microsoft Windows Server 2003 and 2008.
  - OpenSolaris.
  - FreeBSD.
  - NetBSD.
- A user can
  - Load an *EC2* instance with a custom application environment.
  - Manage network's access permissions.
  - Run the image using as many or as few systems as desired.

# EC2 (cont'd)

---

- Import virtual machine (VM) images from the user environment to an instance through *VM import*.
- *EC2* instances boot from an AMI (Amazon Machine Image) digitally signed and stored in S3.
- Users can access:
  - Images provided by Amazon.
  - Customize an image and store it in S3.
- An *EC2* instance is characterized by the resources it provides:
  - VC (Virtual Computers) – virtual systems running the instance.
  - CU (Compute Units) – measure computing power of each system.
  - Memory.
  - I/O capabilities.

# Instance types

- Standard instances: micro (StdM), small (StdS), large (StdL), extra large (StdXL); small is the default.
- High memory instances: high-memory extra large (HmXL), high-memory double extra large (Hm2XL), and high-memory quadruple extra large (Hm4XL).
- High CPU instances: high-CPU extra large (HcpuXL).
- Cluster computing: cluster computing quadruple extra large (Cl4XL).

Instance name	API name	Platform (32/64-bit)	Memory (GB)	Max EC2 compute units	I-memory (GB)	I/O (M/H)
StdM		32 and 64	0.633	1 VC; 2 CUs		
StdS	m1.small	32	1.7	1 VC; 1 CU	160	M
StdL	m1.large	64	7.5	2 VCs; 2 × 2 CUs	85	H
StdXL	m1.xlarge	64	15	4 VCs; 4 × 2 CUs	1,690	H
HmXL	m2.xlarge	64	17.1	2 VCs; 2 × 3.25 CUs	420	M
Hm2XL	m2.2xlarge	64	34.2	4 VCs; 4 × 3.25 CUs	850	H
Hm4XL	m2.4xlarge	64	68.4	8 VCs; 8 × 3.25 CUs	1,690	H
HcpuXL	c1.xlarge	64	7	8 VCs; 8 × 2.5 CUs	1,690	H
Cl4XL	cc1.4xlarge	64	18	33.5 CUs	1,690	H

# S3 – Simple Storage System

---

- Service designed to store large objects; an application can handle an unlimited number of objects ranging in size from 1 byte to 5 TB.
- An object is stored in a bucket and retrieved via a unique, developer-assigned key; a bucket can be stored in a Region selected by the user.
- Supports a minimal set of functions: write, read, and delete; it does not support primitives to copy, to rename, or to move an object from one bucket to another.
- The object names are global.
- S3 maintains for each object: the name, modification time, an access control list, and up to 4 KB of user-defined metadata.



## S3 (cont'd)

---

- Authentication mechanisms ensure that data is kept secure.
- Objects can be made public, and rights can be granted to other users.
- S3 computes the MD5 of every object written and returns it in a field called ETag.
- A user is expected to compute the MD5 of an object stored or written and compare this with the ETag; if the two values do not match, then the object was corrupted during transmission or storage.

# Elastic Block Store (EBS)

---

- Provides persistent block level storage volumes for use with *EC2* instances; suitable for database applications, file systems, and applications using raw data devices.
- A volume appears to an application as a raw, unformatted and reliable physical disk; the range 1 GB -1 TB.
- An *EC2* instance may mount multiple volumes, but a volume cannot be shared among multiple instances.
- EBS supports the creation of snapshots of the volumes attached to an instance and then uses them to restart the instance.
- The volumes are grouped together in Availability Zones and are automatically replicated in each zone.

# SimpleDB

---

- Non-relational data store. Supports store and query functions traditionally provided only by relational databases.
- Supports high performance Web applications; users can store and query data items via Web services requests.
- Creates multiple geographically distributed copies of each data item.
- It manages automatically:
  - The infrastructure provisioning.
  - Hardware and software maintenance.
  - Replication and indexing of data items.
  - Performance tuning.

# SQS - Simple Queue Service

---

- Hosted message queues are accessed through standard SOAP and Query interfaces.
- Supports automated workflows - *EC2* instances can coordinate by sending and receiving SQS messages.
- Applications using SQS can run independently and asynchronously, and do not need to be developed with the same technologies.
- A received message is “locked” during processing; if processing fails, the lock expires and the message is available again.
- Queue sharing can be restricted by IP address and time-of-day.

# CloudWatch

---

- Monitoring infrastructure used by application developers, users, and system administrators to collect and track metrics important for optimizing the performance of applications and for increasing the efficiency of resource utilization.
- Without installing any software a user can monitor either seven or eight pre-selected metrics and then view graphs and statistics for these metrics.
- When launching an Amazon Machine Image (AMI) the user can start the CloudWatch and specify the type of monitoring:
  - Basic Monitoring - free of charge; collects data at five-minute intervals for up to seven metrics.
  - Detailed Monitoring - subject to charge; collects data at one minute interval.

# AWS services introduced in 2012

---

- *Route 53* - low-latency DNS service used to manage user's DNS public records.
- *Elastic MapReduce (EMR)* - supports processing of large amounts of data using a hosted Hadoop running on *EC2*.
- *Simple Workflow Service (SWF)* - supports workflow management; allows scheduling, management of dependencies, and coordination of multiple *EC2* instances.
- *ElastiCache* - enables web applications to retrieve data from a managed in-memory caching system rather than a much slower disk-based database.
- *DynamoDB* - scalable and low-latency fully managed NoSQL database service.

# AWS services introduced in 2012 (cont'd)

---

- *CloudFront* - web service for content delivery.
- *Elastic Load Balancer* - automatically distributes the incoming requests across multiple instances of the application.
- *Elastic Beanstalk* - handles automatically deployment, capacity provisioning, load balancing, auto-scaling, and application monitoring functions.
- *CloudFormation* - allows the creation of a stack describing the infrastructure for an application.

# Elastic Beanstalk

---

- Handles automatically the deployment, capacity provisioning, load balancing, auto-scaling, and monitoring functions.
- Interacts with other services including *EC2*, *S3*, *SNS*, Elastic Load Balance and AutoScaling.
- The management functions provided by the service are:
  - Deploy a new application version (or rollback to a previous version).
  - Access to the results reported by CloudWatch monitoring service.
  - Email notifications when application status changes or application servers are added or removed.
  - Access to server log files without needing to login to the application servers.
- The service is available using: a Java platform, the PHP server-side description language, or the .NET framework.



# SaaS services offered by Google

---

- *Gmail* - hosts Emails on Google servers and provides a web interface to access the Email.
- *Google docs* - a web-based software for building text documents, spreadsheets and presentations.
- *Google Calendar* - a browser-based scheduler; supports multiple user calendars, calendar sharing, event search, display of daily/weekly/monthly views, and so on.
- *Google Groups* - allows users to host discussion forums to create messages online or via Email.
- *Picasa* - a tool to upload, share, and edit images.
- *Google Maps* - web mapping service; offers street maps, a route planner, and an urban business locator for numerous countries around the world

# PaaS services offered by Google

---

- *AppEngine* - a developer platform hosted on the cloud.
  - Initially supported Python, Java was added later.
  - The database for code development can be accessed with GQL (Google Query Language) with a SQL-like syntax.
- *Google Co-op* - allows users to create customized search engines based on a set of facets/categories.
- *Google Drive* - an online service for data storage.
- *Google Base* - allows users to load structured data from different sources to a central repository, a very large, self-describing, semi-structured, heterogeneous database.

# PaaS and SaaS services from Microsoft

---

- *Windows Azure* - an operating system; has 3 components:
  - Compute - provides a computation environment.
  - Storage - for scalable storage.
  - Fabric Controller - deploys, manages, and monitors applications.
- *SQL Azure* - a cloud-based version of the SQL Server.
- *Azure AppFabric*, formerly .NET Services - a collection of services for cloud applications.

# Open-source platforms for private clouds

---

- *Eucalyptus* - can be regarded as an open-source counterpart of Amazon's EC2.
- *Open-Nebula* - a private cloud with users actually logging into the head node to access cloud functions. The system is centralized and its default configuration uses the NFS file system.
- *Nimbus* - a cloud solution for scientific applications based on Globus software; inherits from Globus:
  - The image storage.
  - The credentials for user authentication.
  - The requirement that a running Nimbus process can **ssh** into all compute nodes.

# Eucalyptus

---

- *Virtual Machines* - run under several VMMs including Xen, KVM, and VMware.
- *Node Controller* - runs on server nodes hosting a VM and controls the activities of the node.
- *Cluster Controller* - controls a number of servers.
- *Cloud Controller* - provides the cloud access to end-users, developers, and administrators.
- *Storage Controller* - provides persistent virtual hard drives to applications. It is the correspondent of EBS.
- *Storage Service (Walrus)* - provides persistent storage; similar to S3, it allows users to store objects in buckets.

[EUCALYPTUS](#)
[RIGHTSCALE  
MYCLOUD](#)
[ECOSYSTEM TOOLS](#)
[SECURITY](#)
[DOCUMENTATION](#)
[GET EUCALYPTUS](#)
[FastStart](#)
[Free Trial](#)
[Eucalyptus](#)
[Euca2ools](#)
[Community Cloud](#)
[Source](#)
[Home](#) > [Eucalyptus Cloud](#) > [Get Eucalyptus](#)

## DOWNLOAD EUCALYPTUS

First time using Eucalyptus? Try [Eucalyptus FastStart](#).

### 1. Download and Install Eucalyptus

Choose a distribution:

[CentOS 5](#)
[CentOS 6](#)
[RHEL 5](#)
[RHEL 6](#)
[Ubuntu 10.04 LTS](#)
[Ubuntu 12.04 LTS](#)
[Source](#)
[Versions prior to Eucalyptus 3.1](#)
[Nightlies](#)
[Release Notes](#)

Looking for [Euca2ools](#)?

### 2. Configure Your Cloud

[Documentation](#)
[Engage \(Q&A\)](#)
[Consulting](#)
[Education](#)
[Support](#)

### 3. Use Your Cloud

To help get you started, we have prepared pre-packaged virtual machines ready to run in your Eucalyptus cloud.

[Download images](#)

Or check out a variety of [use cases](#).

### Learn About Eucalyptus For

[MANAGERS](#)
[ARCHITECTS](#)
[APPLICATION ARCHITECTS](#)
[ADMINISTRATORS](#)
[DEVELOPERS](#)
[USERS](#)

### Euca2ools

Eucalyptus supported command-line tools.

[Get Euca2ools](#)

### Ecosystem Tools

Find tools developed for Amazon EC2 and S3 which are compatible with Eucalyptus.

[Get tools](#)

# Cloud storage diversity and vendor lock-in

---

- Risks when a large organization relies on a single cloud service provider:
  - Cloud services may be unavailable for a short or an extended period of time.
  - Permanent data loss in case of a catastrophic system failure.
  - The provider may increase the prices for service.
- Switching to another provider could be very costly due to the large volume of data to be transferred from the old to the new provider.
- A solution is to replicate the data to multiple cloud service providers, similar to data replication in RAID.

# Energy use and ecological impact

---

- The energy consumption of large-scale data centers and their costs for energy and for cooling are significant.
- In 2006, the 6,000 data centers in the U.S consumed  $61 \times 10^9$  KWh of energy, 1.5% of all electricity consumption, at a cost of \$4.5 billion.
- The energy consumed by the data centers was expected to double from 2006 to 2011 and peak instantaneous demand to increase from 7 GW to 12 GW.
- The greenhouse gas emission due to the data centers is estimated to increase from  $116 \times 10^9$  tones of  $\text{CO}_2$  in 2007 to 257 tones in 2020 due to increased consumer demand.
- The effort to reduce energy use is focused on computing, networking, and storage activities of a data center.



# Energy use and ecological impact (cont'd)

---

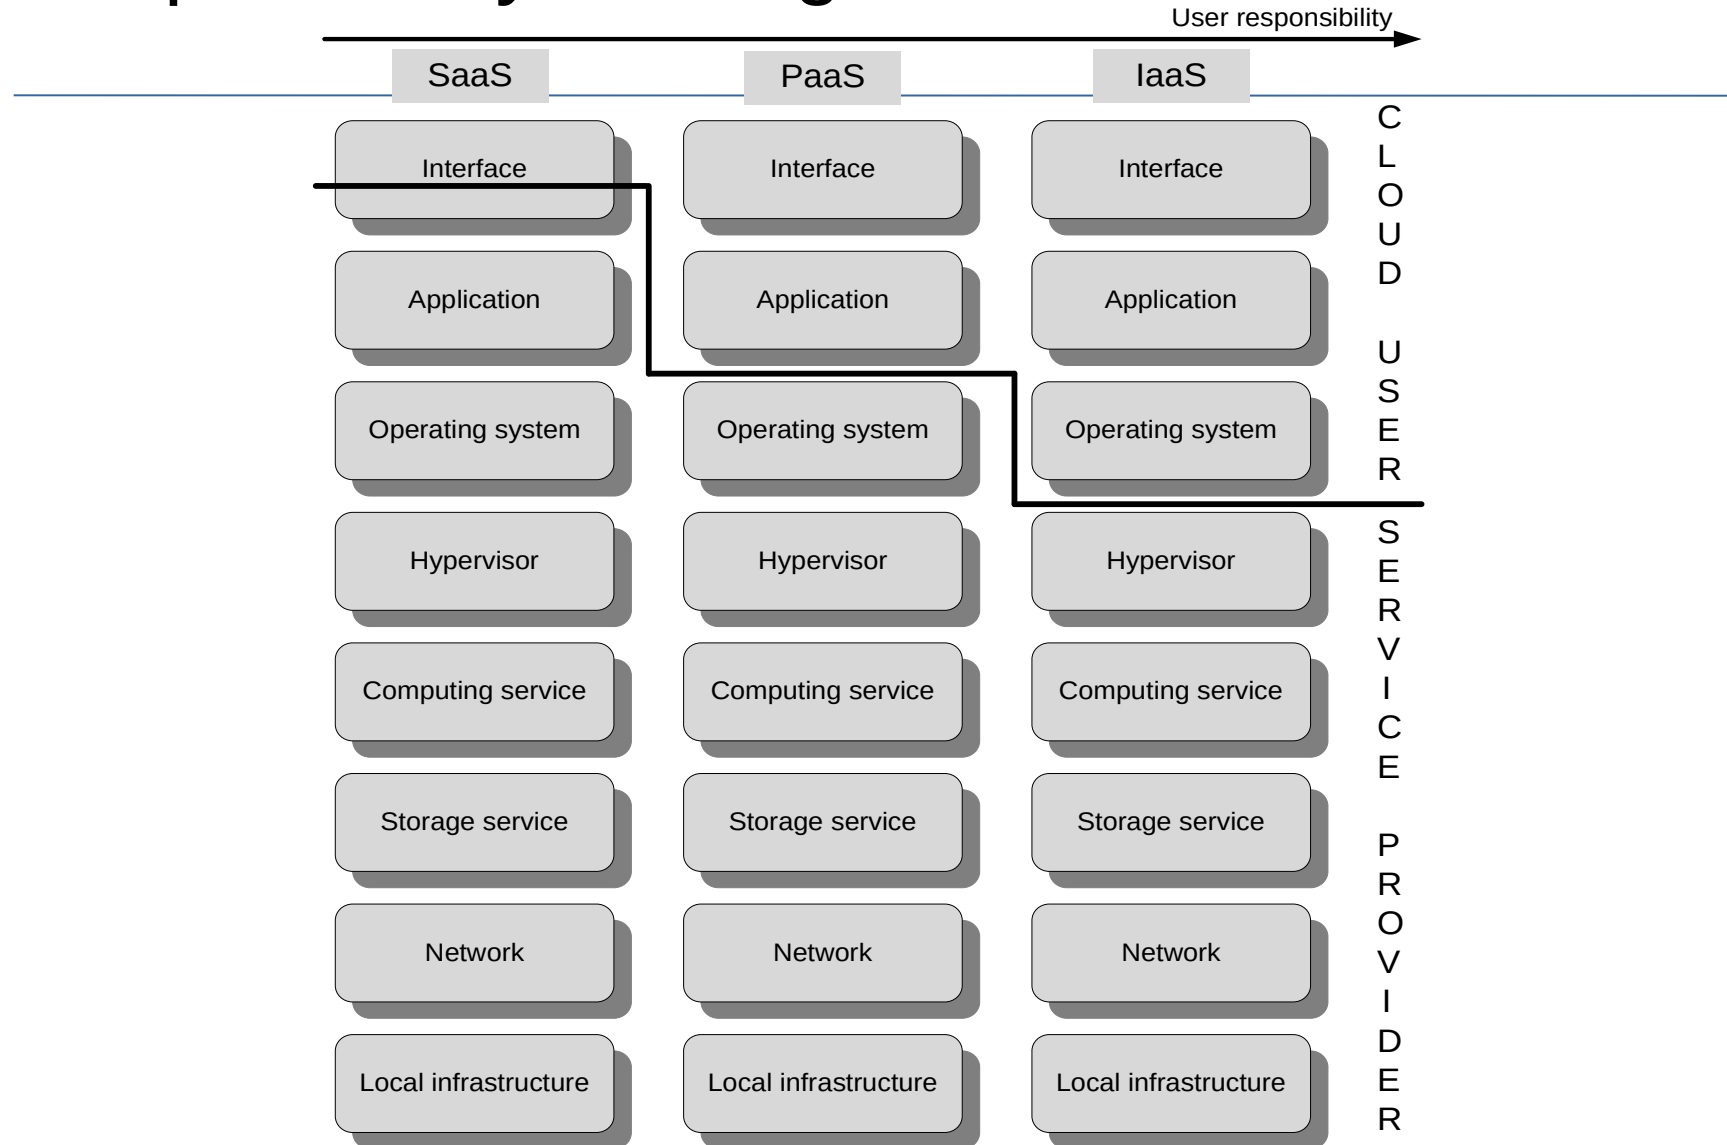
- Operating efficiency of a system is captured by the *performance per Watt of power*.
- The performance of supercomputers has increased 3.5 times faster than their operating efficiency – 7,000% versus 2,000% during the period 1998 – 2007.
- A typical Google cluster spends most of its time within the 10-50% CPU utilization range; there is a mismatch between server workload profile and server energy efficiency.

# Service Level Agreement (SLA)

---

- SLA - a negotiated contract between the customer and CSP; can be legally binding or informal. Objectives:
  - Identify and define the customer's needs and constraints including the level of resources, security, timing, and QoS.
  - Provide a framework for understanding; a critical aspect of this framework is a clear definition of classes of service and the costs.
  - Simplify complex issues; clarify the boundaries between the responsibilities of clients and CSP in case of failures.
  - Reduce areas of conflict.
  - Encourage dialog in the event of disputes.
  - Eliminate unrealistic expectations.
- Specifies the services that the customer receives, rather than how the cloud service provider delivers the services.

# Responsibility sharing between user and CSP



# User security concerns

---

- Potential loss of control/ownership of data.
- Data integration, privacy enforcement, data encryption.
- Data remanence after de-provisioning.
- Multi tenant data isolation.
- Data location requirements within national borders.
- Hypervisor security.
- Audit data integrity protection.
- Verification of subscriber policies through provider controls.
- Certification/Accreditation requirements for a given cloud service.

# Credits, references and reading material

---

- *Cloud Computing: Theory and Practice*

Dan C. Marinescu

Chapter 3

- *Distributed and Cloud Computing*

K. Hwang, G. Fox and J. Dongarra

Morgan Kaufmann, 2012